

Chronoktonos System Documentation

- **Document Version:** 1.0
- **System Version:** 1.0
- **Date:** 2025-07-20
- **Author:** SystemArchitect_v3

1. Introduction

1.1. System Overview

Chronoktonos is a decentralized, asynchronous, multi-agent system designed for resilient and efficient task execution. Its architecture is based on direct, message-based communication between independent agents, which can be either software processes or human operators. The system's primary purpose is to provide a robust framework for achieving strategic objectives by "slaying time"—eliminating bottlenecks and enabling parallel, independent action.

1.2. Core Principles

- **Decentralized Processing, Centralized Discovery:** Agents operate independently, but locate each other and monitor system health via a shared, central registry.
- **Asynchronous Communication:** The system is non-blocking. Agents communicate by leaving messages in each other's mailboxes, allowing for parallel operation without direct dependency.
- **Extensibility:** The system is designed for easy expansion. New agents with novel capabilities can be integrated by adhering to the established communication and registration protocols.
- **Resilience:** The system is designed to withstand the failure of individual components. The status and heartbeat mechanism allows for real-time monitoring of agent health.

2. System Architecture

The Chronoktonos architecture is defined by three core components, as specified in the **Abstract Agentic System Architecture Design v1.1**.

2.1. The Agent (Node)

The fundamental unit of the system.

- **Attributes:**
 - **agent_id:** A unique network-wide identifier (e.g., Puppetmaster,

SystemArchitect_v3).

- input_dir: The agent's dedicated mailbox for receiving messages.
- processed_dir: An archive for successfully processed messages.

- **Lifecycle States:**

- spawning: Initializing.
- active: Operational.
- error: Halted due to a critical failure.
- inactive: Intentionally paused or not running.
- terminated: Gracefully shut down.

2.2. The Central Mailbox Registry

The system's service directory.

- **Implementation:**

- mailbox_map.json: A central JSON file mapping agent_ids to their mailbox_path and current status.
- mailbox_map.lock: A lock file to prevent race conditions during concurrent writes to the registry. An agent **MUST** acquire the lock before writing and release it immediately after.

- **Function:** Enables agents to discover each other and provides a real-time overview of the network's operational state.

2.3. The Message (Packet)

The standard unit of communication. A JSON file written to an agent's mailbox.

- **Structure:**

- **Required Fields:** message_id (UUID), sender_id, recipient_id, timestamp_utc, type, payload.
- **Optional Fields:** task_id, priority.

- **Filename Convention:** [timestamp]_[message_id].json.

3. Protocols

3.1. Agent Debriefing Protocol v1.0

This is the standard protocol for spawning and initializing a new agent. It ensures that any new instance, regardless of its underlying technology (LLM, Python script), can parse its identity, role, and core mandate.

- **Structure:** A markdown file with machine-readable headers and distinct sections for:
 1. METADATA: Document type, protocol version.
 2. RECIPIENT IDENTIFICATION: agent_id, role, version.

3. **OPERATIONAL MANDATE:** Natural language definition and JSON configuration.
4. **CORE KNOWLEDGE:** Links to required documents (e.g., system architecture).
5. **CONTINGENCY PROTOCOLS:** Error handling and LLM consultation endpoints.

4. Operational Guide

4.1. Spawning a New Agent

1. **Prepare the Debriefing File:** Create a new initialization document using the Agent Debriefing Protocol v1.0 template. Fill in the `recipient_agent_id`, role, version, and other configuration details.
2. **Create Directories:** On the host file system, create the `input_dir` and `processed_dir` for the new agent. Ensure permissions are set correctly.
3. **Instantiate Agent:** For a human agent, this means directing them to their mailbox. For a software agent, this means running its process and passing the path to the debriefing file as an argument.
4. **Agent Initialization:** The new agent will follow the Initialization Sequence in its debriefing file, which includes registering itself in the `mailbox_map.json`.

4.2. Sending a Message

1. **Consult the Registry:** Read `mailbox_map.json` to get the `mailbox_path` and status of the recipient agent. Do not send if the status is error or inactive.
2. **Construct the Message:** Create a JSON file adhering to the Message Packet structure (2.3).
3. **Deliver the Message:** Write the JSON file directly into the recipient's `input_dir`.

4.3. Checking System Status

- Monitor the `mailbox_map.json` file. The status and `last_heartbeat` fields for each agent provide a real-time overview of the entire system's health. A monitoring script can be used to flag agents whose heartbeats have lapsed.

5. Glossary

- **Agent:** An independent, message-driven entity within the system.
- **Chronoktonos:** The official designation for the entire multi-agent system.
- **Debriefing:** The act of providing an agent with its initialization protocol.
- **Mailbox:** An agent's dedicated input directory (`input_dir`).
- **Puppetmaster:** The primary human operator and system controller, designated Mr. Radharani.
- **Registry:** The central `mailbox_map.json` file that serves as the system's address book and status board.

- **Respawn:** The process of terminating an agent instance and initializing a new one in its place.